



Homework 03

NAME:

STUDENT ID:

Numpy Introduction

1a) Create two numpy arrays (a and b). a should be all integers between 25-34 (inclusive), and b should be ten evenly spaced numbers between 1-6. Print all the results below:

- i) Cube (i.e. raise to the power of 3) all the elements in both arrays (element-wise)
- ii) Add both the cubed arrays (e.g., $[1,2] + [3,4] = [4,6]$)
- iii) Sum the elements with even indices of the added array.
- iv) Take the square root of the added array (element-wise square root)___

In [1]:

```
import numpy as np
a = np.arange(25,35)
b = np.linspace(1,6,10)
print(a**3)
print(b**3)
print(a**3+b**3)
print(sum([(a**3+b**3)[i] for i in range(10) if i%2==0]))
print(np.sqrt(a**3+b**3))
```

```
[15625 17576 19683 21952 24389 27000 29791 32768 35937 39304]
[  1.          3.76406036  9.40877915 18.96296296 33.45541838
 53.91495199 81.37037037 116.85048011 161.38408779 216.]
]
[ 15626.          17579.76406036 19692.40877915 21970.96296296
 24422.45541838 27053.91495199 29872.37037037 32884.85048011
 36098.38408779 39520.]
]
125711.618656
[ 125.00399994 132.58870261 140.32964327 148.22605359 156.27685503
 164.48074341 172.83625306 181.34180566 189.99574755 198.7963782
 4]
```

1b) Append b to a, reshape the appended array so that it is a 4x5, 2d array and store the results in a variable called m. Print m.

In [2]:

```
m = np.concatenate([a, b]).reshape(4,5)
print(m)
```

```
[ [ 25.          26.          27.          28.          29.          ]
  [ 30.          31.          32.          33.          34.          ]
  [  1.          1.55555556   2.11111111   2.66666667   3.22222222 ]
  [  3.77777778   4.33333333   4.88888889   5.44444444   6.          ]]
```

1c) Extract the third and the fourth column of the m matrix. Store the resulting 4x2 matrix in a new variable called m2. Print m2.

In [3]:

```
m2 = m[:,2:4]
print(m2)
```

```
[ [ 27.          28.          ]
  [ 32.          33.          ]
  [  2.11111111   2.66666667 ]
  [  4.88888889   5.44444444 ]]
```

1d) Take the dot product of m2 and m store the results in a matrix called m3. Print m3. Note that Dot product of two matrices $A.B = A^T B$

In [4]:

```
m3 = np.dot(m2.T,m)
print(m3)
```

```
[ [ 1655.58024691  1718.4691358   1781.35802469  1844.24691358
    1907.13580247]
  [ 1713.2345679   1778.74074074  1844.24691358  1909.75308642
    1975.25925926 ]]
```

1e) Round the m3 matrix to three decimal points. Store the result in place and print the new m3.

In [5]:

```
m3 = np.round(m3,decimals=3)
print(m3)
```

```
[ [ 1655.58   1718.469  1781.358  1844.247  1907.136]
  [ 1713.235   1778.741  1844.247  1909.753  1975.259 ]]
```

1f) Sort the m3 array so that the highest value is at the bottom right and the lowest value is at the top left. Print the sorted m3 array.

In [6]:

```
print(np.sort(m3,axis=None).reshape(2,5))
```

```
[ [ 1655.58   1713.235  1718.469  1778.741  1781.358]
  [ 1844.247  1844.247  1907.136  1909.753  1975.259 ]]
```

NumPy and Masks

2a) create an array called 'f' where the values are cosine(x) for x from 0 to pi with 50 equally spaced values in f

- print f
- use a 'mask' and print an array that is True when $f \geq 1/2$ and False when $f < 1/2$
- create and print an array sequence that has only those values where $f \geq 1/2$

In [7]:

```
f = np.cos(np.linspace(0,np.pi,50))
print(f)
mask = f>=1/2
print(mask)
print(f[mask])
```

```
[ 1.          0.99794539  0.99179001  0.98155916  0.96729486  0.949055
75      0.92691676  0.90096887  0.8713187   0.8380881   0.80141362  0.761445
96      0.71834935  0.67230089  0.6234898   0.57211666  0.51839257  0.462538
29      0.40478334  0.34536505  0.28452759  0.22252093  0.1595999   0.096023
03      0.03205158 -0.03205158 -0.09602303 -0.1595999   -0.22252093 -0.284527
59     -0.34536505 -0.40478334 -0.46253829 -0.51839257 -0.57211666 -0.623489
8     -0.67230089 -0.71834935 -0.76144596 -0.80141362 -0.8380881   -0.871318
7     -0.90096887 -0.92691676 -0.94905575 -0.96729486 -0.98155916 -0.991790
01     -0.99794539 -1.          ]
[ True  True  True  True  True  True  True  True  True  True  True  Tr
ue    True  True  True  True  True False False False False False False Fal
se    False False False False False False False False False False False Fal
se    False False False False False False False False False False False Fal
se    False False]
```

```
[ 1.          0.99794539  0.99179001  0.98155916  0.96729486  0.949055
75      0.92691676  0.90096887  0.8713187   0.8380881   0.80141362  0.761445
96      0.71834935  0.67230089  0.6234898   0.57211666  0.51839257]
```

NumPy and 2 Variable Prediction

Let 'x' be the number of miles a person drives per day and 'y' be the dollars spent on buying car fuel (per day).

We have created 2 numpy arrays each of size 100 that represent x and y.
x (number of miles) ranges from 1 to 10 with a uniform noise of (0,1/2)
y (money spent in dollars) will be from 1 to 20 with a uniform noise (0,1)

In [8]:

```
# seed the random number generator with a fixed value
import numpy as np
np.random.seed(500)

x=np.linspace(1,10,100)+ np.random.uniform(low=0,high=.5,size=100)
y=np.linspace(1,20,100)+ np.random.uniform(low=0,high=1,size=100)
print ('x = ',x)
print ('y= ',y)
```

```
x = [ 1.34683976  1.12176759  1.51512398  1.55233174  1.40619168
 1.65075498  1.79399331  1.80243817  1.89844195  2.00100023
 2.3344038  2.22424872  2.24914511  2.36268477  2.49808849
 2.8212704  2.68452475  2.68229427  3.09511169  2.95703884
 3.09047742  3.2544361  3.41541904  3.40886375  3.50672677
 3.74960644  3.64861355  3.7721462  3.56368566  4.01092701
 4.15630694  4.06088549  4.02517179  4.25169402  4.15897504
 4.26835333  4.32520644  4.48563164  4.78490721  4.84614839
 4.96698768  5.18754259  5.29582013  5.32097781  5.0674106
 5.47601124  5.46852704  5.64537452  5.49642807  5.89755027
 5.68548923  5.76276141  5.94613234  6.18135713  5.96522091
 6.0275473  6.54290191  6.4991329  6.74003765  6.81809807
 6.50611821  6.91538752  7.01250925  6.89905417  7.31314433
 7.20472297  7.1043621  7.48199528  7.58957227  7.61744354
 7.6991707  7.85436822  8.03510784  7.80787781  8.22410224
 7.99366248  8.40581097  8.28913792  8.45971515  8.54227144
 8.6906456  8.61856507  8.83489887  8.66309658  8.94837987
 9.20890222  8.9614749  8.92608294  9.13231416  9.55889896
 9.61488451  9.54252979  9.42015491  9.90952569  10.00659591
 10.02504265  10.07330937  9.93489915  10.0892334  10.36509991]
y= [ 1.6635012  2.0214592  2.10816052  2.26016496  1.96287558
 2.9554635  3.02881887  3.33565296  2.75465779  3.4250107
 3.39670148  3.39377767  3.78503343  4.38293049  4.32963586
 4.03925039  4.73691868  4.30098399  4.8416329  4.78175957
 4.99765787  5.31746817  5.76844671  5.93723749  5.72811642
 6.70973615  6.68143367  6.57482731  7.17737603  7.54863252
 7.30221419  7.3202573  7.78023884  7.91133365  8.2765417
 8.69203281  8.78219865  8.45897546  8.89094715  8.81719921
 8.87106971  9.66192562  9.4020625  9.85990783  9.60359778
 10.07386266  10.6957995  10.66721916  11.18256285  10.57431836
 11.46744716  10.94398916  11.26445259  12.09754828  12.11988037
 12.121557  12.17613693  12.43750193  13.00912372  12.86407194
 13.24640866  12.76120085  13.11723062  14.07841099  14.19821707
 14.27289001  14.30624942  14.63060835  14.2770918  15.0744923
 14.45261619  15.11897313  15.2378667  15.27203124  15.32491892
 16.01095271  15.71250558  16.29488506  16.70618934  16.56555394
 16.42379457  17.18144744  17.13813976  17.69613625  17.37763019
 17.90942839  17.90343733  18.01951169  18.35727914  18.16841269
 18.61813748  18.66062754  18.81217983  19.44995194  19.7213867
 19.71966726  19.78961904  19.64385088  20.69719809  20.07974319]
```

3a) Find Expected value of x and the expected value of y

In [9]:

```
print(np.mean(x))
print(np.mean(y))
```

```
5.78253254159
11.0129816833
```

3b) Find variance of distributions of x and y

In [10]:

```
print(np.var(x))
```

```
7.03332752948
```

In [11]:

```
print(np.var(y))
```

```
30.1139035755
```

3c) Find co-variance of x and y.

In [12]:

```
print(np.mean(x*y) - np.mean(x)*np.mean(y))
```

```
14.5111663945
```

3d) Assuming that number of dollars spent in car fuel is only dependant on the miles driven, by a linear relationship.

Write code that uses a linear predictor to calculate a predicted value of y for each x ie $y_{\text{predicted}} = f(x) = y_0 + mx$.

In [13]:

```
X = np.column_stack([np.ones(len(x)), x])
W = np.linalg.inv(X.T.dot(X)).dot(X.T).dot(y)
print(W)
```

```
[-0.9175436    2.06320072]
```

In [14]:

```
y_predicted = np.dot(X,W)
print(y_predicted)
```

```
[ 1.86125717  1.39688809  2.20846128  2.28522836  1.98371207
 2.48829527  2.78382468  2.80124813  2.9993232  3.21092152  3.
8988
 3.67152796  3.7228942  3.9571493  4.23651436  4.9033035
 4.62116978  4.61656787  5.46829307  5.18342105  5.45873164
 5.79701128  6.12915141  6.11562653  6.31753758  6.81864709
 6.61027849  6.86515115  6.43505522  7.35780389  7.65775187
 7.46087825  7.38719373  7.85455455  7.66325667  7.88892606
 8.00622544  8.33721481  8.95468038  9.08103323  9.33034895
 9.78539799 10.00879629 10.06070164  9.53754157 10.38056671
10.36512531 10.72999716 10.42269073 11.25028634 10.81276185
10.97218988 11.35052091 11.83583685 11.38990445 11.51849632
12.58177632 12.49147206 12.98850691 13.14956122 12.50588416
13.35028889 13.5506705  13.31658991 14.17094102 13.947246
13.74018137 14.51931443 14.74126735 14.79877137 14.96739089
15.28759454 15.66049665 15.1916755  16.05043004 15.57498655
16.42533161 16.18461169 16.53654675 16.70687695 17.01300263
16.86428603 17.31062607 16.95616347 17.54476017 18.08227006
17.57177784 17.49875711 17.92425351 18.80438359 18.91989301
18.77061069 18.51812677 19.5277969  19.72807224 19.76613158
19.8657155  19.58014745 19.89856998 20.46773797]
```

3e) Predict y for each value in x, put the error into an array called y_error

In [15]:

```
y_error = y_predicted - y
```

In [16]:

```
print(y_error)
```

```
[ 0.19775597 -0.62457111  0.10030076  0.02506341  0.02083649 -0.467168
23 -0.24499418 -0.53440482  0.24466541 -0.21408918  0.50209852  0.277750
29 -0.06213923 -0.42578118 -0.0931215   0.86405311 -0.1157489   0.315583
88  0.62666017  0.40166149  0.46107377  0.47954311  0.3607047   0.178389
04  0.58942116  0.10891094 -0.07115518  0.29032384 -0.74232081 -0.190828
63  0.35553767  0.14062095 -0.39304511 -0.0567791  -0.61328502 -0.803106
76 -0.77597321 -0.12176065  0.06373323  0.26383402  0.45927925  0.123472
38  0.60673379  0.20079382 -0.0660562   0.30670405 -0.33067419  0.062778
-0.75987212  0.67596798 -0.65468531  0.02820071  0.08606832 -0.261711
43 -0.72997592 -0.60306068  0.40563939  0.05397013 -0.02061681  0.285489
28 -0.7405245   0.58908804  0.43343988 -0.76182107 -0.02727604 -0.325644
01 -0.56606805 -0.11129392  0.46417555 -0.27572093  0.5147747   0.168621
42  0.42262995 -0.08035574  0.72551112 -0.43596616  0.71282602 -0.110273
37 -0.16964259  0.14132301  0.58920807 -0.31716141  0.17248631 -0.739972
78  0.16712997  0.17284167 -0.33165948 -0.52075457 -0.43302563  0.635970
9  0.30175553  0.10998314 -0.29405306  0.07784496  0.00668554  0.046464
31  0.07609646 -0.06370343 -0.79862812  0.38799477]
```

3f) Write code that calculates the root mean square error(RMSE), that is root of average of y-error squared

In [17]:

```
RMSE = np.sqrt(np.mean(y_error**2))
```

In [18]:

```
RMSE
```

Out[18]:

```
0.41767772366856104
```